

LEAD	Returns an expression from a later row that is a defined offset from the current row.
FIRST_VALUE	Returns NULL if no row at specified position.
LAST_VALUE	Returns the first value in the current window frame. Requires window ordering to be meaningful.

```
--what is the current balance of the account per a given row..this query did not had good performance
before sql2012
select * from Transactions
where AccountID = 29825

select *
, SUM(Amount) over (
    partition by AccountID
    order by TransactionDate, TransactionID
    rows between unbounded preceding and current row
) as FinalBalance
from Transactions
where AccountID = 29825
order by AccountID , TransactionDate, TransactionID

--i want to 1 step backwards in the set
select *
, lag(Amount,1,0) over (partition by AccountID order by TransactionDate, TransactionID
rows between unbounded preceding and current row
) as FinalBalance
from Transactions
where AccountID = 29825
order by AccountID , TransactionDate, TransactionID
```

What is pivoting?

Pivoting data is rotating data from a rows-based orientation to a columns-based orientation

Distinct values from a single column are projected across as headings for other columns - may include aggregation

Category	Qty	Orderyear	Category	2006	2007	2008
Dairy Products	12	2006	Beverages	1842	3996	3694
Grains/Cereals	10	2006	Condiments	962	2895	1441
Dairy Products	5	2006	Confections	1357	4137	2412
Produce	9	2006	Dairy Products	2086	4374	2689
Produce	40	2006	Grains/Cereals	549	2636	1377
Seafood	10	2006	Meat/Poultry	950	2189	1060
Produce	35	2006	Produce	549	1583	858
Condiments	15	2006	Seafood	1286	3679	2716
Grains/Cereals	6	2006				
Grains/Cereals	15	2006				
Condiments	20	2006				
Confections	40	2006				
Dairy Products	25	2006				
Dairy Products	40	2006				
Dairy Products	20	2006				

Pivoted data

Pivoting includes three phases:

1. Grouping determines which element gets a row in the result set
2. Spreading provides the distinct values to be pivoted across
3. Aggregation performs an aggregation function (such as SUM)

These 3 phases for both T-SQL standard and SQL standard remain the same for both T-SQL standard and SQL standard.

SELECT Category, [2006], [2007], [2008]
FROM (SELECT Category, Qty, Orderyear
FROM Sales.CategoryQtyYear) AS D
PIVOT(SUM(QTY) FOR orderyear
IN([2006],[2007],[2008]))
AS pvt;

```
SELECT vendorID, [250] AS Emp1, [251] AS Emp2, [256] AS Emp3, [257] AS Emp4, [260] AS Emp5
FROM
```

```
(SELECT PurchaseOrderID, EmployeeID, VendorID
FROM Purchasing.PurchaseOrderHeader) p
PIVOT
(
  COUNT (PurchaseOrderID)
FOR EmployeeID IN
( [250], [251], [256], [257], [260] )
) AS pvt
ORDER BY pvt.VendorID;
```

VendorID	Emp1	Emp2	Emp3	Emp4	Emp5
1492	2	5	4	4	4
1494	2	5	4	5	4
1496	2	4	4	5	5

Note that u generally don't use pivot in sql as it mainly in used in excel or reporting services. So pivoting is generally done up in the application layer.

```
CREATE TABLE pvt (VendorID int, Emp1 int, Emp2 int,
                  Emp3 int, Emp4 int, Emp5 int);
GO
INSERT INTO pvt VALUES (1,4,3,5,4,4);
INSERT INTO pvt VALUES (2,4,1,5,5,5);
INSERT INTO pvt VALUES (3,4,3,5,4,4);
GO
```

```
SELECT VendorID, Employee, Orders
FROM
  (SELECT VendorID, Emp1, Emp2, Emp3, Emp4, Emp5
   FROM pvt) p
UNPIVOT
  (Orders FOR Employee IN
    (Emp1, Emp2, Emp3, Emp4, Emp5)
  )AS unpvt;
GO
```

VendorID	Employee	Orders
1	Emp1	4
1	Emp2	3

Writing queries with grouping sets:

GROUPING SETS subclause builds on T-SQL GROUP BY clause

Allows multiple groupings to be defined in same query

Alternative to use of UNION ALL to combine multiple outputs (each with different GROUP BY) into one result set

```
SELECT <column list with aggregate(s)>
  FROM <source>
  GROUP BY
  GROUPING SETS(
```

```
  (<column_name>),--one or more columns
  (<column_name>),--one or more columns
  () -- empty parentheses if aggregating all rows
);
```

```
SELECT TerritoryID, CustomerID, SUM(TotalDue) AS TotalAmountDue
  FROM Sales.SalesOrderHeader
 GROUP BY
  GROUPING SETS((TerritoryID),(CustomerID),());
```

TerritoryID	CustomerID	TotalAmountDue
NULL	30116	211671.2674
NULL	30117	919801.8188
NULL	30118	313671.5352
NULL	NULL	123216786.1159

3	NULL	8913299.2473
6	NULL	18398929.188
9	NULL	11814376.0952
1	NULL	18061660.371
7	NULL	8119749.346

--lets say we want to get total sales by customer and territory

```

--WHAT IS THE TOTAL SALES IN TERRITORY INDEPENDENT OF THE CUSTOMER
SELECT
C.CustomerID
,C.TerritoryID
,SUM(SOD.LineTotal) AS TOTALSALES
FROM Sales.Customer AS C
INNER JOIN Sales.SalesOrderHeader AS SOH ON SOH.CustomerID = C.CustomerID
INNER JOIN Sales.SalesOrderDetail AS SOD ON SOD.SalesOrderID = SOH.SalesOrderID
GROUP BY
C.CustomerID
,C.TerritoryID
UNION ALL
SELECT
NULL
,NULL
,SUM(SOD.LineTotal) AS TOTALSALES
FROM Sales.Customer AS C
INNER JOIN Sales.SalesOrderHeader AS SOH ON SOH.CustomerID = C.CustomerID
INNER JOIN Sales.SalesOrderDetail AS SOD ON SOD.SalesOrderID = SOH.SalesOrderID
ORDER BY CustomerID

-- now we are writting a lot of code. instead use grouping sets...EASY AND WE CAN ADD MORE GROUPING SETS
EASILY
SELECT
C.CustomerID
,C.TerritoryID
,SUM(SOD.LineTotal) AS TOTALSALES
FROM Sales.Customer AS C
INNER JOIN Sales.SalesOrderHeader AS SOH ON SOH.CustomerID = C.CustomerID
INNER JOIN Sales.SalesOrderDetail AS SOD ON SOD.SalesOrderID = SOH.SalesOrderID
GROUP BY GROUPING SETS(
(C.CustomerID,C.TerritoryID)
,((C.CustomerID,C.TerritoryID)
,()
)
)
ORDER BY CustomerID, TerritoryID

```

CUBE and ROLLUP:

CUBE provides shortcut for defining grouping sets given a list of columns

All possible combinations of grouping sets are created

```

SELECT TerritoryID, CustomerID, SUM(TotalDue) AS TotalAmountDue
FROM Sales.SalesOrderHeader
GROUP BY ROLLUP(TerritoryID, CustomerID)

```

ROLLUP provides shortcut for defining grouping sets, creates combinations assuming input columns form a hierarchy

```

SELECT TerritoryID, CustomerID, SUM(TotalDue) AS TotalAmountDue
FROM Sales.SalesOrderHeader
GROUP BY ROLLUP(TerritoryID, CustomerID)

```

ORDER BY TerritoryID, CustomerID;

DML statements (INSERT, UPDATE, and DELETE):

Using INSERT to add data:

The INSERT...VALUES statement inserts a single row by default
INSERT INTO Production.UnitMeasure (Name, UnitMeasureCode, ModifiedDate)
VALUES ('N'Square Yards', N'Y2', GETDATE());
GO

Not Standard → Values can be used as a select & combination

Table and row constructors add multi-row capability to INSERT...VALUES
INSERT INTO Production.UnitMeasure (Name, UnitMeasureCode, ModifiedDate)
VALUES
(N'Square Feet', N'F2', GETDATE()),
(N'Square Inches', N'I2', GETDATE());

Using INSERT with DEFAULT constraint:

DEFAULT constraints are used to assign a value to a column when none is specified in the INSERT statement
Defaults are defined in CREATE or ALTER TABLE statement

INSERT can omit columns which have defaults defined

Applies to IDENTITY and NULL too

VALUES clause can use DEFAULT keyword

If no default constraint assigned, NULL inserted

If you added a field called Country with a DEFAULT constraint of 'USA' you could INSERT data using the following command

INSERT INTO Production.UnitMeasure (Name, UnitMeasureCode, ModifiedDate, Country)
VALUES ('N'Square Miles', N'M2', GETDATE(), DEFAULT);

Using INSERT with SELECT and EXEC:

INSERT...SELECT is used to insert the result set of a query into an existing table
INSERT INTO Production.UnitMeasure (Name, UnitMeasureCode, ModifiedDate)
SELECT Name, UnitMeasureCode, ModifiedDate
FROM Sales.TempUnitTable
WHERE ModifiedDate < '20080101';

INSERT...EXEC is used to insert the result of a stored procedure or dynamic SQL expression into an existing table
INSERT INTO Production.UnitMeasure (Name, UnitMeasureCode, ModifiedDate)
EXEC Production.Temp_UOM
@numrows = 5, @catid=1;

Not Standard ↗
BULK INSERT data
There is a file
Statement from

Using SELECT INTO:

SELECT...INTO is similar to INSERT...SELECT but SELECT...INTO creates a new table each time the statement is executed

Copies column names, data types, and nullability of identity property.

Does not copy constraints or indexes

① triggers ② permissions

SELECT Name, UnitMeasureCode, ModifiedDate
INTO Production.TempUOMTable
FROM Production.UnitMeasure
WHERE orderdate < '20080101';

Heap is a table without a Heap.

Using IDENTITY:

IDENTITY property of a column generates sequential numbers automatically for insertion into a table

Can specify optional set seed and increment values

Only one column in a table may have IDENTITY property defined

IDENTITY column is omitted in INSERT statements

Functions provided to retrieve last generated value

Creates a table with using the IDENTITY property with a starting number of 100 and incremented by 10 as each row is added

CREATE TABLE Production.IdentityProducts(
productid int IDENTITY(100,10) NOT NULL,
productname nvarchar(40) NOT NULL,
categoryid int NOT NULL,
unitprice money NOT NULL)

② identity selects last identity generated from Production.IdentityProducts

TRUNCATE MERGE

Standard

Using SEQUENCES:
Sequence objects new in SQL Server 2012
Independent objects in database

More flexible than the IDENTITY property
Can be used as default value for a column
Manage with CREATE/ALTER/DROP statements

Retrieve value with the NEXT VALUE FOR clause
-- Define a sequence

CREATE SEQUENCE dbo.InvoicesSeq AS INT START WITH 5 INCREMENT BY 5;

-- Retrieve next available value from sequence

SELECT NEXT VALUE FOR dbo.InvoicesSeq;

Using UPDATE to modify data:

Updates all rows in a table or view

Set can be filtered with a WHERE clause

Set can be defined with a JOIN clause

Only columns specified in the SET clause are modified

Updates the ModifiedDate using a the GETDATE function for the record that has 'M2' in the UnitMeasureCode

UPDATE Production.UnitMeasure

SET ModifiedDate = (GETDATE())

WHERE UnitMeasureCode = 'M2'.

If no WHERE clause is specified, all records in the Production.UnitMeasure will be updated

Using MERGE to modify data:

MERGE modifies data based on one of the following conditions

When the source matches the target

When the target has no match in the source

MERGE INTO schema_name.table_name AS TargetTbl

USING (SELECT <select_list>) AS SourceTbl

ON (TargetTbl.col1 = SourceTbl.col1)

WHEN MATCHED THEN

UPDATE SET col2 = SourceTbl.col2

WHEN NOT MATCHED THEN

INSERT (<column_list>)

VALUES (<value_list>);

--SELECT SOME DATA INTO A NEW TABLE

--SELECT

--ProductID

--,Name

--INTO PRODUCTS

--FROM Production.Product;

--SELECT
*
FROM PRODUCTS
WHERE NAME LIKE 'A%',
FOR XML PATH

--RETURN AS XML
SELECT
*
FROM PRODUCTS
WHERE NAME LIKE 'A%',
FOR XML PATH('PRODUCT'), ROOT('PRODUCTS')

--RETURN AS XML...CHANGE NAMES
SELECT

IDENT_CURRENT('dbo.tbl') returns last identity value for a table. So in case of an error, this value might not be found stored in the table.
Still not correct
Identity property
does not enforce uniqueness.

Set col1 = col2, col2 = col
As noted earlier, we can run DML against views. Now view is a table expression. In general, DML can be run against all table expressions (CTE, UDF, derived tables, view).
MERGE alternative to UPDATE based on a JOIN.

with WHEN MATCHED clause is a standard

PRODUCTID AS "@id"--WE HAVE TO USE DOUBLE INVERTED COMMAS AS @ HAS SPECIAL MEANING IN SQL...AND @ IS USED TO SPECIFY THAT
--WE WANT THIS TO BE A ATTRIBUTE RATHER THAN AN ELEMENT

```
NAME AS "@name"
FROM PRODUCTS
WHERE NAME LIKE 'A%'
```

FOR XML PATH('PRODUCT'), ROOT('PRODUCTS')

--UPDATES USING XML...KEEP THE CASING TO LOWER CASE WHEN USING XML
DECLARE @xml XML = N'

```
<PRODUCTS>
<PRODUCT id="1" name="Adjustable Race!" />
<PRODUCT id="879" name="All-Purpose Bike Stand!" />
<PRODUCT id="712" name="AWC Logo Cap!" />
<PRODUCT name="Ramneek!" />
</PRODUCTS>;
```

XML is case-sensitive

```
SELECT
XT.XC.value('@id', 'INT') AS PRODUCTID
,XT.XC.value('@name', 'NVARCHAR(1000)') AS NAME
FROM @xml.nodes('/PRODUCTS/PRODUCT') AS XT(XC);
```

WITH SRC AS (

```
SELECT
XT.XC.value('@id', 'INT') AS PRODUCTID
,XT.XC.value('@name', 'NVARCHAR(1000)') AS NAME
FROM @xml.nodes('/PRODUCTS/PRODUCT') AS XT(XC)
```

)

-- NOW MERGE THIS DATA INTO THE PRODUCTS TABLE
MERGE INTO PRODUCTS AS DEST

USING SRC ON SRC.PRODUCTID = DEST.PRODUCTID

WHEN NOT MATCHED THEN

INSERT (NAME) VALUES(SRC.NAME)

WHEN MATCHED THEN

UPDATE SET NAME = SRC.NAME;

SELECT

PRODUCTID AS "@id"--WE HAVE TO USE DOUBLE INVERTED COMMAS AS @ HAS SPECIAL MEANING IN SQL...AND @ IS IS USED TO SPECIFY THAT
--WE WANT THIS TO BE A ATTRIBUTE RATHER THAN AN ELEMENT

```
,NAME AS "@name"
FROM PRODUCTS
WHERE NAME LIKE 'A%' OR NAME LIKE 'Ram%'
FOR XML PATH('PRODUCT'), ROOT('PRODUCTS')
```

--ALSO MERGE FOR DELETES.

DECLARE @xml XML = N'

<PRODUCTS>

```
<PRODUCT id="1" name="Adjustable Race!" />
<PRODUCT id="879" name="All-Purpose Bike Stand!" />
<PRODUCT id="712" delete="true" name="AWC Logo Cap!" />
<PRODUCT id="1000" name="Ramneek!" />
</PRODUCTS>;
```

WITH SRC AS (

```
SELECT
XT.XC.value('@id', 'INT') AS PRODUCTID
,XT.XC.value('@name', 'NVARCHAR(1000)') AS NAME
,ISNULL(XT.XC.value('@delete', 'BIT'),0) AS DODELETE
FROM @xml.nodes('/PRODUCTS/PRODUCT') AS XT(XC)
```

--NOW MERGE THIS DATA INTO THE PRODUCTS TABLE
MERGE INTO PRODUCTS AS DEST

USING SRC ON SRC.PRODUCTID = DEST.PRODUCTID
WHEN NOT MATCHED THEN

INSERT (NAME) VALUES(SRC.NAME)

WHEN MATCHED AND SRC.DODELETE=0 THEN
UPDATE SET NAME = SRC.NAME

WHEN MATCHED AND src.DODELETE=1 THEN
delete;

```
SELECT
    PRODUCTID AS "@id"--WE HAVE TO USE DOUBLE INVERTED COMMAS AS @ HAS SPECIAL MEANING IN SQL...AND @ IS USED TO SPECIFY THAT
    --WE WANT THIS TO BE A ATTRIBUTE RATHER THAN AN ELEMENT
    ,NAME AS "@name"
FROM PRODUCTS
WHERE NAME LIKE 'A%' OR NAME LIKE 'Ram%'
FOR XML PATH('PRODUCT'), ROOT('PRODUCTS')
```

Using DELETE to remove data:

DELETE operates on a set

Set may be filtered with a WHERE clause

Deletion of each row is logged in database's transaction log

DELETE may be rolled back if statement issued within a user-defined transaction or if an error is encountered

```
DELETE FROM Production.UnitMeasure
WHERE UnitMeasureCode = 'Y2';
```

If no WHERE clause is specified, all records in the Production.UnitMeasure will be deleted

Using TRUNCATE TABLE to remove all data:

TRUNCATE TABLE clears the entire table

Storage is physically deallocated, rows not individually removed

Minimally logged

Can be rolled back if TRUNCATE issued within a transaction

```
TRUNCATE TABLE Production.UnitMeasure
```

Note: if you have a lot of rows, then doing a 'delete from' without a where clause will take a long time. Then use the 'truncate' option

PRIMARY KEY constraint:

A PRIMARY KEY is an important concept of designing a database table as it provides an attribute or set of attributes used to uniquely identify each row in the table

A table can only have one primary key which is created using a primary key constraint and enforced by creating a unique index on the primary key columns. So underscores are not only used to speed up queries by avoiding full table scans

To add a PRIMARY KEY constraint to an existing table use the following command

```
ALTER TABLE Production.TransactionHistoryArchive
ADD CONSTRAINT PK_TransactionHistoryArchive_TransactionID
PRIMARY KEY CLUSTERED (TransactionID);
```

Primary key) for key). Unique key) constraint.

Truncate resets the identity
Seed.
the table is empty!

FOREIGN KEY constraint:

A FOREIGN KEY is a column or combination of columns that are used to establish a link between data in two tables. The columns used to create the primary key in one table are also used to create the foreign key constraint and can be used to reference data in the same table or in another table

A foreign key does not have to reference a primary key, it can be defined to reference a unique constraint in either the same table or in another table. ~~NULL is allowed in for. key cols. unless restricted otherwise~~ unique

To add a FOREIGN KEY constraint to an existing table use the following command

```
ALTER TABLE Sales.SalesOrderHeadersSalesReason
ADD CONSTRAINT FK_SalesReason
```

FOREIGN KEY (SalesReasonID)

```
REFERENCES Sales.SalesReason (SalesReasonID)
```

```
ON DELETE CASCADE
```

```
ON UPDATE CASCADE;
```

~~UNIQUE constraints~~

~~instead of cascade~~ ~~use cascade relationship or even in same table~~

key of one table is in another table

Creating a UNIQUE constraint automatically creates a corresponding unique index

No COUNT no. of rows
that will be affected reported.
you can use if you want should not be affected reported.

A table can have one primary key but multiple unique key constraints.

To create a UNIQUE constraint while creating a table use the following command

```
CREATE TABLE Production.TransactionHistoryArchive4  
(TransactionID int NOT NULL,  
CONSTRAINT AK_TransactionID UNIQUE(TransactionID) );
```

CHECK constraints:

A CHECK constraint is created in a table to specify the data values that are acceptable in one or more columns

```
ALTER TABLE DBO.NewTable  
ADD ZipCode int NULL  
CONSTRAINT CHK_ZipCode  
CHECK (ZipCode LIKE '[0-9][0-9][0-9][0-9][0-9]');
```

```
But then then that col can have only one NULL value
```

To create a CHECK constraint after creating a table use the following command

```
ALTER TABLE Sales.CountryRegionCurrency  
ADD CONSTRAINT Default_Country  
DEFAULT 'USA' FOR CountryRegionCode
```

DEFAULT constraints:

A DEFAULT constraint is a special case of a column default that is applied when an INSERT statement doesn't explicitly assign a particular value. In other words, the column default is what the column will get as a value by default

To create a DEFAULT constraint on an existing table use the following command

```
ALTER TABLE Sales.CountryRegionCurrency  
ADD CONSTRAINT Default_Country  
DEFAULT 'USA' FOR CountryRegionCode
```

DML triggers:

A DML trigger is a special type of stored procedure that automatically fires when any valid DML event takes place regardless of whether or not any rows are affected

AFTER and INSTEAD OF triggers

DML triggers can include complex T-SQL statements used to enforce business rules and provide data integrity when an

INSERT, UPDATE, or DELETE command is executed

INSTEAD OF Can also be created on Views

The following trigger will prints a message to the client when anyone tries to add or change data in the Customer table

CREATE TRIGGER reminder1 ON Sales.Customer

Trigger considered part of the transaction

AFTER INSERT, UPDATE

AS RAISERROR ('Notify Customer Relations', 16, 10);

then includes the event that caused the trigger to fire

Trigger to fire from action

Note: we might want to use the new throw statement instead of the raiseerror used here

Note: if the rules are too complex to be enforced by constraints, then use trigger

Note: other use could be to update data in some denormalized tables. But if performance is hampered, then you might want to use some sort of background job instead of a trigger to update the denormalized tables.

Trigger fired even if no rows affected by DML statements

OUTPUT clause:

The OUTPUT clause is used to return information from, or expressions based on, each row affected by an INSERT, UPDATE, DELETE, or MERGE statement. These results can be returned to the processing application for use in such things as confirmation messages or archiving

The following example deletes all rows in the ShoppingCartItem table. The clause OUTPUT deleted.* specifies that all

columns in the deleted rows, be returned to the calling application which in this case was the Query Editor

```
DELETE FROM Sales.ShoppingCartItem OUTPUT DELETED.* WHERE ShoppingCartID = 20621;
```

--Verify the rows in the table matching the WHERE clause have been deleted.

```
SELECT COUNT(*) AS [Rows in Table]
```

```
FROM Sales.ShoppingCartItem
```

```
WHERE ShoppingCartID = 20621;
```

Trigger fired even if no rows affected by DML statements

```
DECLARE @tmp TABLE (ProductID INT PRIMARY KEY);
```

UPDATE Production.Product SET

```
Name = UPPER(Name)
```

```
OUTPUT INSERTED.ProductID INTO @tmp
```

```
WHERE ListPrice > 10;
```

```
SELECT * FROM @tmp;
```

T-SQL batches:

T-SQL batches are collections of one or more T-SQL statements sent to SQL Server as a unit for parsing, optimization, and execution and are terminated with the GO clause

Batches are also boundaries for the scope of a variable

Output info to ??

Can also be done in a table or view

Output - before update

Output - before insert

Some statements (e.g., CREATE FUNCTION, CREATE PROCEDURE, CREATE VIEW) may not be combined with others in the same batch

```
CREATE VIEW HumanResources.EmployeeList
AS
SELECT BusinessEntityID, JobTitle, HireDate, VacationHours
FROM HumanResources.Employee;
GO
```

Batches are parsed for syntax as a unit

Syntax errors cause the entire batch to be rejected

Runtime errors may allow the batch to continue after failure, by default

Batches can contain error-handling code

```
--Valid batch
INSERT INTO Production.UnitMeasure (Name, UnitMeasureCode, ModifiedDate)
VALUES (N'Square Footage', N'F4', GETDATE());
(N'Square Inches', N'I2', GETDATE());
GO
--Invalid batch
INSERT INTO dbo.t1 VALUE(1,2,N'abc');
INSERT INTO dbo.t1 VALUES(2,3,N'def');
GO
```

T-SQL Variables:

Variables are objects that allow storage of a value for use later in the same batch

Variables are defined with the DECLARE keyword and begin with @

Beginning with SQL Server 2008, variables can be declared and initialized in the same statement

Variables are always local to the batch in which they're declared and go out of scope when the batch ends

--Declare, initialize, and use a variable

```
DECLARE @SalesPerson_id INT = 5;
SELECT OrderYear, COUNT(DISTINCT CustomerID) AS CustCount
FROM (
    SELECT YEAR(OrderDate) AS OrderYear, CustomerID
    FROM Sales.SalesOrderHeader
    WHERE SalesPersonID = @SalesPerson_id
) AS DerivedYear
GROUP BY OrderYear;
```

Working with Variables:

Values can be assigned with a SET command or a SELECT statement

SET can only assign one variable at a time. SELECT can assign multiple variables at a time

When using SELECT to assign a value, make sure that exactly one row is returned by the query

--Declare and initialize variables

```
DECLARE @numrows INT = 3, @catid INT = 2;
--Use variables to pass parameters to procedure
EXEC Production.ProdsByCategory
    @numrows = @numrows, @catid = @catid;
GO
```

Working with Synonyms:

A synonym is an alias or link to an object stored either on the same SQL Server instance or on a linked server

Synonyms can point to tables, views, procedures, and functions

Synonyms can be used for referencing remote objects as though they were located locally, or for providing alternative names to other local objects

Use the CREATE, ALTER, and DROP commands to manage synonyms

```
--Create a synonym for the Product table in AdventureWorks CREATE SYNONYM dbo.MyProduct
FOR AdventureWorks.Production.Product;
GO
-- Query the Product table by using the synonym.
SELECT ProductID, Name
FROM MyProduct
WHERE ProductID < 5;
GO
```

T-SQL control-of-flow language:

SQL Server provides additional language elements that control the flow and execution of T-SQL statements in batches, stored procedures, and multi-statement functions

Control-of-flow elements allow you to specify statements need to be performed in a specified order or not at all. The default is for statements to execute sequentially, however you can use IF...ELSE, BEGIN...END, WHILE, RETURN, and others to control the flow of your batch files or stored procedures

IF OBJECT_ID ('Production.Product', 'U') IS NOT NULL

PRINT 'I am here and contain data, so don't delete me'

IF...ELSE:
IF...ELSE uses a predicate to determine the flow of the code

The code in the IF block is executed if the predicate evaluates to TRUE

The code in the ELSE block is executed if the predicate evaluates to FALSE or UNKNOWN

Very useful when combined with the EXISTS operator

IF OBJECT_ID ('Production.Product', 'U') IS NOT NULL

PRINT 'I am here and contain data, so don't delete me'

ELSE

PRINT 'Table not found, so feel free to create one'

GO

WHILE:

WHILE enables statements to execute in the WHILE block as long as the predicate evaluates to TRUE and doesn't stop executing until the predicate evaluates to FALSE or UNKNOWN

Execution can be altered by BREAK or CONTINUE

```
DECLARE @BusinessEntID AS INT = 1, @Title AS NVARCHAR(50);
WHILE @BusinessEntID <=10
BEGIN
    SELECT @Title = JobTitle FROM HumanResources.Employee
    WHERE BusinessEntityID = @BusinessEntID;
    PRINT @Title;
    SET @BusinessEntID += 1;
END;
```

Structured exception handling:

Structured exception handling allows a centralized response to runtime errors

TRY to run a block of commands and CATCH any errors

Execution moves to the CATCH block of commands when an error occurs

No need to check every statement to see if an error occurred

If error you decide whether the transaction should be rolled back, errors logged, etc.

Not all errors can be caught by TRY / CATCH:

Syntax or compile errors

Some name resolution errors

Querying the ERROR object:

Common ERROR object properties and ERROR object functions:

Property	Function to Query	Description
Number	ERROR_NUMBER	Unique number assigned to the error
Message	ERROR_MESSAGE	Error message text
Severity	ERROR_SEVERITY	Severity class (1-25)
Procedure Name	ERROR_PROCEDURE	Name of the procedure or trigger that raised the error
Line Number	ERROR_LINE	Number of the line that raised the error in the batch, procedure, trigger, or function

Values returned correspond to sys.messages view

TRY and CATCH blocks:

TRY block defined by BEGIN TRY...END TRY statements

Place all code that might raise an error between them