

No code may be placed between END TRY and BEGIN CATCH

TRY and CATCH blocks may be nested

CATCH block defined by BEGIN CATCH...END CATCH

Execution moves to the CATCH block when catchable errors occur within the TRY block

-- Generate a divide-by-zero error.

```
SELECT 1/0;
END TRY
BEGIN CATCH
    SELECT
        ERROR_NUMBER() AS ErrorNumber
        ,ERROR_SEVERITY() AS ErrorSeverity
        ,ERROR_STATE() AS ErrorCode
        ,ERROR_PROCEDURE() AS ErrorProcedure
        ,ERROR_LINE() AS ErrorLine
        ,ERROR_MESSAGE() AS ErrorMessage;
END CATCH;
GO
```

ErrorNumber	ErrorSeverity	ErrorState	ErrorProcedure	ErrorLine	ErrorMessage
8134	16	1	NULL	3	Divide by zero error encountered.

Not all errors can be caught by TRY / CATCH:

Syntax, compilation errors, and name resolution errors

```
BEGIN TRY
    -- Table does not exist; object name resolution
    -- error not caught.
```

```
SELECT * FROM IDontExist;
```

```
END TRY
BEGIN CATCH
    SELECT
        ERROR_NUMBER() AS ErrorNumber
        ,ERROR_MESSAGE() AS ErrorMessage;
END CATCH
```

Invalid object name 'IDontExist'.

THROW statement:

SQL Server 2012 provides the new THROW statement

Successor to the RAISERROR statement

Does not require defining errors in the sys.messages table

THROW allows choices when handling errors:

Handle specific errors in the local CATCH block

Pass errors to another process

Use THROW:

With parameters to pass a user-defined error

Without parameters to re-raise the original error (must be within a CATCH block)

```
BEGIN TRY
    SELECT 100/0 AS 'Problem';
END TRY
BEGIN CATCH
    PRINT 'Code inside CATCH is beginning'
    PRINT 'MyError: ' + CAST(ERROR_NUMBER()
        AS VARCHAR(255));
    THROW;
END CATCH
```

```
(0 row(s) affected)
Code inside CATCH is beginning
MyError: 8134
Msg 8134, Level 16, State 1, Line 6
Divide by zero error encountered.
```

Transactions:

A transaction is a group of tasks defined as a unit of work that must succeed or fail together – no partial completion is permitted

```
-- Two tasks that make up a unit of work
INSERT INTO Sales.SalesOrderHeader...
INSERT INTO Sales.SalesOrderDetail...
```

SQL Server uses locking mechanisms and the transaction log to support transactions

The need for transactions: issues with batches:

Some runtime errors during a batch may result in unacceptable partial success:

Part of the batch succeeds and part fails, leaving behind the results from the part of the batch that succeeded
--Batch without transaction management

```
BEGIN TRY
    INSERT INTO Sales.SalesOrderHeader... --Insert succeeds
    INSERT INTO Sales.SalesOrderDetail... --Insert fails
END TRY
BEGIN CATCH
    --First row still in Sales.SalesOrderHeader Table
    SELECT ERROR_NUMBER()
    ...
END CATCH;
```

Select if you are in a transaction
to see to see

Transactions extend batches:

Transaction commands identify blocks of code that must succeed or fail together and provide points where database engine can roll back, or undo, operations:

```
BEGIN TRY
    BEGIN TRANSACTION
        INSERT INTO Sales.SalesOrderHeader... --Succeeds
        INSERT INTO Sales.SalesOrderDetail... --Fails
    COMMIT TRANSACTION -- If no errors, transaction completes
END TRY
BEGIN CATCH
    --Inserted rows still exist in Sales.SalesOrderHeader
    ROLLBACK TRANSACTION --Any transaction work undone
END CATCH;
```

'ACID' properties
of transactions

BEGIN TRANSACTION marks the starting point of an explicit, user-defined transaction
Transactions last until a COMMIT statement is issued, a ROLLBACK is manually issued, or the connection is broken and the system issues a ROLLBACK
Transactions are local to a connection and cannot span connections
In your T-SQL code: Mark the start of the transaction's work

```
BEGIN TRY
    BEGIN TRANSACTION -- marks beginning of transaction
    INSERT INTO Sales.SalesOrderHeader... --Completed
    INSERT INTO Sales.SalesOrderDetail... --Completed
    ...
    COMMIT ensures all of the transaction's modifications are made a permanent part of the database
    COMMIT frees resources, such as locks, used by the transaction
```

In your T-SQL code: If a transaction is successful, commit it

```
BEGIN TRY
    BEGIN TRAN -- marks beginning of transaction
    INSERT INTO Sales.SalesOrderHeader...
    INSERT INTO Sales.SalesOrderDetail...
    COMMIT TRAN -- mark the transaction as complete
END TRY
```

A ROLLBACK statement undoes all modifications made in the transaction by reverting the data to the state it was in at the beginning of the transaction

ROLLBACK frees resources, such as locks, held by the transaction

Before rolling back, you can test the state of the transaction with the XACT_STATE function

→ In a nested TRAN level rollback
COMMIT at each COMMIT actually
only the last changes of entire
commits to db.

In your T-SQL code: If an error occurs, ROLLBACK to the point of the BEGIN TRANSACTION statement

```
BEGIN CATCH  
    SELECT ERROR_NUMBER() -- sample error handling  
    ROLLBACK TRAN  
END CATCH;
```

SQL Server does not automatically roll back transactions when errors occur

To roll back, either use ROLLBACK statements in error-handling logic or enable XACT_ABORT

XACT_ABORT specifies whether SQL Server automatically rolls back the current transaction when a runtime error occurs

When SET XACT_ABORT is ON, the entire transaction is terminated and rolled back on error, unless occurring in TRY block

SET XACT_ABORT OFF is the default setting

Change XACT_ABORT value with the SET command:

```
SET XACT_ABORT ON;
```

System catalog views:

Built-in views that provide information about the system catalog

Use standard query methods to return metadata

Column lists, JOIN, WHERE, ORDER BY

Some views are filtered to display only user objects, some views include system objects

--Pre-filtered to exclude system objects

```
SELECT name, object_id, schema_id, type, type_desc  
FROM sys.tables;
```

--Includes system and user objects

```
SELECT name, object_id, schema_id, type, type_desc  
FROM sys.objects;
```

AUTO COMMIT  *ndo not do a 'BEGIN TRAN'*
Single DML Statement  *but do ROLLBACK at COMMIT*

Information schema views:

Views stored in the INFORMATION_SCHEMA system schema

Return system metadata per ISO standard, used by third-party tools

Maps standard names (catalog, domain) to SQL Server names (database, user-defined data type)

```
SELECT TABLE_CATALOG, TABLE_SCHEMA,  
       TABLE_NAME, TABLE_TYPE  
  FROM INFORMATION_SCHEMA.TABLES;
```

```
SELECT VIEW_CATALOG, VIEW_SCHEMA, VIEW_NAME,  
      TABLE_CATALOG, TABLE_SCHEMA, TABLE_NAME, COLUMN_NAME  
  FROM INFORMATION_SCHEMA.VIEW_COLUMN_USAGE;
```

```
SELECT VIEW_CATALOG, VIEW_SCHEMA, VIEW_NAME,  
      TABLE_CATALOG, TABLE_SCHEMA, TABLE_NAME, COLUMN_NAME  
  FROM INFORMATION_SCHEMA.VIEW_COLUMN_USAGE  
 WHERE COLUMN_NAME = 'BusinessEntityID';
```

System metadata functions:

Return information about settings, values, and objects in SQL Server

Come in a variety of formats

Some marked with a @@ prefix, sometimes incorrectly referred to as global variables: @@VERSION

Some marked with a () suffix, similar to arithmetic or string functions: ERROR_NUMBER()

Some special functions marked with a \$ prefix: \$PARTITION

Queried with a standard SELECT statement:

```
SELECT @@VERSION AS SQL_Version;  
SELECT SERVERPROPERTY('ProductVersion') AS version;  
SELECT SERVERPROPERTY('Collation') AS collation;
```

Querying DMVs and functions:

Dynamic management views are queried like standard views:

```
SELECT session_id, login_time, program_name  
  FROM sys.dm_exec_sessions  
 WHERE is_user_process = 1;
```

Dynamic management functions are queried as table-valued functions, including parameters:

```
SELECT referencing_schema_name, referencing_entity_name, referencing_class_desc  
  FROM sys.dm_sql_referencing_entities(  
      'Sales.SalesOrderHeader', 'OBJECT');
```

About dynamic management objects:

The nearly 200 dynamic management views (DMVs) and functions return server state information
DMVs include catalog information as well as administrative status information, such as object dependencies
DMVs are server-scoped (instance-level) or database-scoped

Requires VIEW SERVER STATE or VIEW DATABASE STATE permission to query DMVs

Underlying structures change over time, so avoid writing SELECT * queries against DMVs

Categories include;

Naming pattern	Description
db	Database-related information
exec	Query execution-related information
io	I/O statistics
os	SQL Server Operating System (SQLoS) information
tran	Transaction-related information

Stored procedures:

Use the EXECUTE or EXEC command before the name of the stored procedure

Pass parameters by position or name, separated by commas when applicable

--no parameters so lists all database

EXEC sys.sp_databases;

--single parameter of name of table

EXEC sys.sp_help N'Sales.Customer';

--multiple named parameters

EXEC sys.sp_tables
@table_name = '%',
@table_owner = N'Sales';

Common system stored procedures:

Database engine procedures can provide general metadata

sp_help, sp_helplanguage

sp_who, sp_lock

Catalog procedures can be used as an alternative to system catalog views and functions:

Name	Description
sp_databases	Lists databases in an instance of SQL Server
sp_tables	Returns a list of tables or views, except synonyms
sp_columns	Returns column information for the specified objects

Unlike system views, there is no option to select which columns to return

Executing system stored procedures:

System stored procedures:

Marked with an sp_ prefix

Stored in a hidden resource database

Logically appear in the sys schema of every user and system database

Best practices for execution include:

Always use EXEC or EXECUTE rather than just calling by name

Include the sys schema name when executing

Name each parameter and specify its appropriate data type

--This example uses EXEC, includes the sys schema name, --and passes the table name as a named Unicode parameter --to a procedure accepting an NVARCHAR(776)
--input parameter.

EXEC sys.sp_help @objname = N'Sales.Customer';

Creating procedures that return rows:

Stored procedures can be wrappers for simple or complex SELECT statements

Procedures may include input and output parameters as well as return values

Use CREATE PROCEDURE statement:

```
CREATE PROCEDURE <schema_name>.proc_name>
(<parameter_list>)
AS
SELECT <body of SELECT statement>;
```

Change procedure with ALTER PROCEDURE statement

No need to drop, recreate

Creating procedures that accept parameters:

Input parameters passed to procedure logically behave like local variables within procedure code

Assign name with @prefix, data type in procedure header

Refer to parameter in body of procedure

```
CREATE PROCEDURE Production.ProdsByProductLine
(@numrows AS int, @ProdLine AS nchar)
```

```
AS
SELECT TOP (@numrows) ProductID,
```

Name, ListPrice

FROM Production.Product

WHERE ProductLine = @ProdLine;

```
--Retrieve top 50 products with product line = M
EXEC Production.ProdsByProductLine 50, 'M'
```

NOTE: SET NOCOUNT ON

Nocount on so that after each statement we don't return number of rows affected back to the client. Reduces the data sent back to the client.

Writing well-performing queries:

Only retrieve what you need

In the SELECT clause, only use needed columns – avoid *

Use a WHERE clause, filter to return only needed rows

Improve search performance of WHERE clause

Avoid expressions that manipulate columns in the predicate

Minimize use of temporary tables or table variables

Use windowing functions or other set-based operations when possible

Avoid cursors and other iterative approaches

Work with your DBA to arrange good indexes to support filters, joins, and ordering

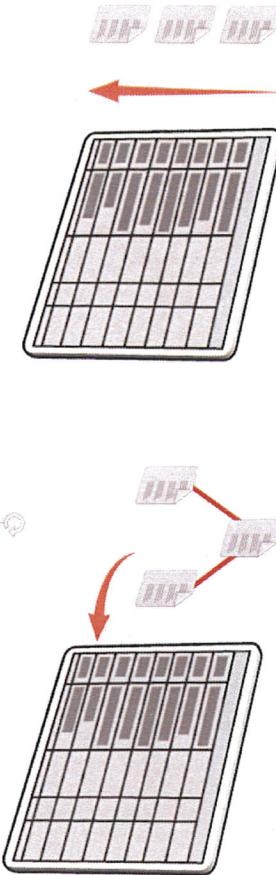
Learn how to address tasks with different query approaches to compare performance

Indexing in SQL Server:

SQL Server accesses data by using indexes or by scanning all rows in a table

Indexes also supports ordering operations such as grouping, joining, and ORDER BY clauses

Table scan: SQL Server reads all table rows



Index seek/scan: SQL Server uses indexes to find rows

SQL Server indexes: performance considerations:

Check query execution plans to see if indexes are present and being used as expected

For query writers who are not DBAs or database developers, the ability to recognize problems with indexes, such as the use of table scans when you expect an index to be used, can be very helpful in tuning an application

Note: while indexes improve the searching of data, there is a slight hit while inserting data.

Distribution statistics:

Distribution statistics describe the distribution and the uniqueness, or selectivity, of data
Statistics, by default, are created and updated automatically
Statistics are used by the query optimizer to estimate the selectivity of data, including the size of the results
Large variances between estimated and actual values might indicate a problem with the estimates, which may be addressed through updating statistics

Avoiding cursors:

Cursors contradict the relational model, which operates on sets
Cursors typically require more code than set-based approach

Cursors typically incur more overhead during execution than a comparable set-based operation
Alternatives to cursors:

Windowing functions

Aggregate functions

Appropriate uses for cursors:

Generating dynamic SQL code

Performing administrative tasks

What is an execution plan?

Review of the process of executing a query:

Parse, resolve, optimize, execute

An execution plan includes information on which tables to access, which indexes, what joins to perform

If statistics exist for a relevant column or index, then the optimizer will use them in its calculations

SQL Server tools provide access to execution plans to show how a query was executed or how it would be executed

Plans available in text format (deprecated), XML format, and graphical renderings of XML

Plan viewer accessible in results pane of SSMS

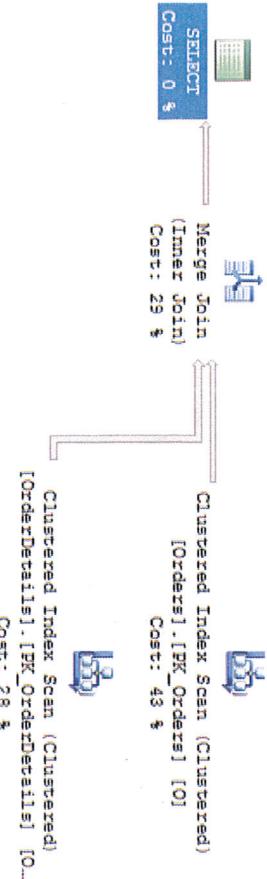
Actual and estimated execution plans:

Execution plans graphically represent the methods that SQL Server uses to execute the statements in a T-SQL query

SSMS provides access to two forms of execution plans:

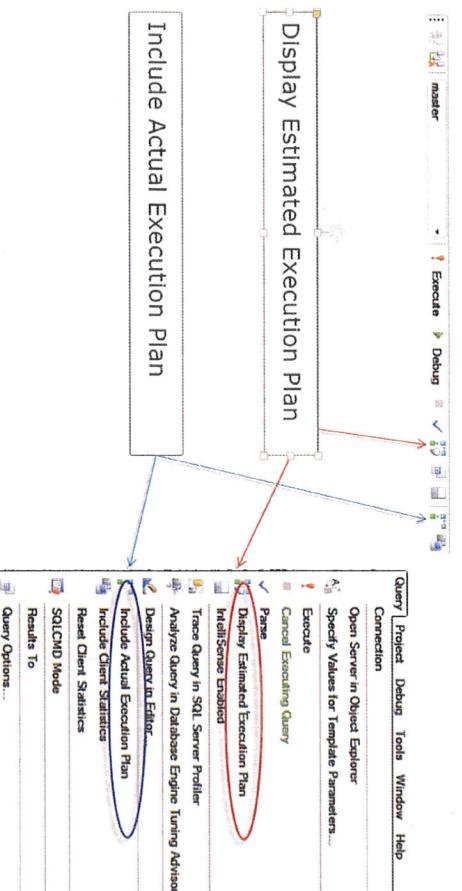
Estimated execution plans do not execute the query. Instead, they display the plan that SQL Server would likely use if the query were run.

Actual execution plans are returned the next time the query is executed. They display the plan that was actually used by SQL Server



Viewing graphical execution plan:

Enable execution plan viewers in SSMS



Interpreting the execution plan:

Read the plan right to left, top to bottom

Hover the mouse pointer over an item to see additional information

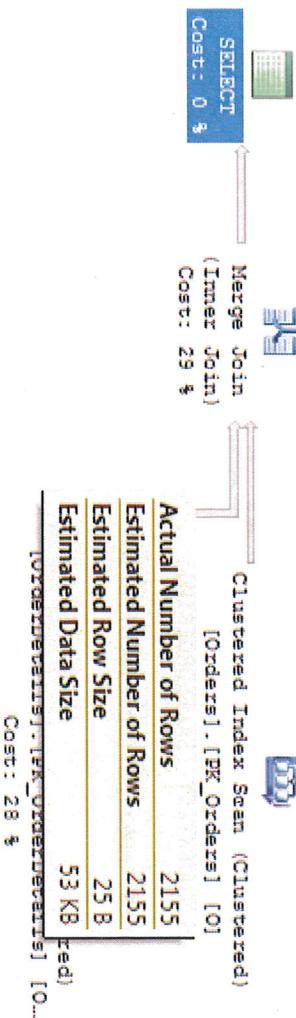
Percentages indicate cost of operator relative to total query

Thickness of lines between operators indicates relative number of rows passing through

For issues, look for thick lines leading into high-cost operators

In an actual execution plan, note any differences between estimated and actual values

Large variances may indicate problems with estimates



Displaying Query Statistics:

SQL Server provides detailed runtime information about the execution of a query

STATISTICS TIME will show time spent parsing and compiling a query

SET STATISTICS TIME ON;

STATISTICS IO will show amount of disk activity generated by a query

SET STATISTICS IO ON;

Query performance:

SET STATISTICS TIME ON;

SET STATISTICS IO ON;

```
--WITHOUT ANY INDEX, A TABLE IS CALLED A HEAP. IT DOES NOT HAVE A CLUSTERED INDEX ON THE TABLE
--IN SQL SERVER 2014, A TABLE CAN BE OF 3 TYPES WHILE IN SQL SERVER 2012 IT CAN BE OF 2 TYPES(HEAP OR
CLUSTERED INDEX)
--SELECT * 
--INTO SALES.INDEXTEST
--FROM SALES.SalesOrderDetail;

EXEC sp_helpindex 'SALES.INDEXTEST';
```

--lets say we want this: by order we want the sum of unitprice*orderqty, groupby sale orderid for a certain specialoffer
--include the actual execution plan and run the query

```
SELECT
    SalesOrderID
    ,SUM(UnitPrice*OrderQty) AS TOTAL
    FROM SALES.INDEXTEST
    WHERE SpecialOfferID=2
    GROUP BY SalesOrderID
    ORDER BY TOTAL;
```

--1499 pages,table scan,hash aggr. :(

--has match aggregate is the grouping ..for grouping either sql server can do a sort or can do a hash operation..typically sorting should be used
--for aggregates. so when u see a hash match in a execution plan, it means u don't have right indexing in place.

--the execution plan also contains a hint that u probably need to create a nonclustered index. You can even generate the indexing code by right clicking the execution plan..now run the query again after switching on the stats 'SET STATISTICS IO ON;' because time stats depend on the locks that happened during the execution and how busy the server is. Now messages tell me per table how much work we did..logical reads tell me how many pages i had to read to get the data --physical reads suggest for how many reads did i have to go the disk to fetch(v/s in-memory reads if the data is already in memory)

--clustering index means u r sorting the table and adding a balanced tree structure on top of the table
CREATE CLUSTERED INDEX CluIdx ON SALES.INDEXTEST (SpecialOfferID);

--and then again run the query. it actually does more page reads..as the table is very big and restructuring it actually caused it to use more space
--1525 pages,clustered index seek,hash aggr. :(but for query for specialofferid=2 has a reduction in the number of page reads and is down to 49

--but we still have a hash aggr.

--if a table already has a clustered index, we have to create non-clustered indexes. take all the cols in the non-clust. index, create a new table with those cols and shovel all the data for those cols into the new table and then add a pointer back to original table --sometime the query analyzer does not use the index that it should have used. then (generally only for testing purposes), you can force the query to use a particular index
DROP INDEX CluIdx ON SALES.INDEXTEST;
CREATE CLUSTERED INDEX CluIdx ON SALES.INDEXTEST (SALESORDERID);

CREATE NONCLUSTERED INDEX NCIdx ON SALES.INDEXTEST(SpecialOfferID);

SELECT

SalesOrderID

,**SUM**(UnitPrice*OrderQty) AS TOTAL
FROM SALES.INDEXTEST **WITH(INDEX(NCIdx))**

WHERE SpecialOfferID=2

GROUP BY SalesOrderID

ORDER BY TOTAL;

--10517 pages,non clustered index seek + table lookup,hash aggr. :(((but for query for specialofferid=2 has a reduction in the number of page reads and is down to 49

--to avoid lookups, include additional cols in the nonclustered index that are being used in query
DROP INDEX CluIdx ON SALES.INDEXTEST;
DROP INDEX NCIdx ON SALES.INDEXTEST;
CREATE CLUSTERED INDEX CluIdx ON SALES.INDEXTEST (SALESORDERID);
CREATE NONCLUSTERED INDEX NCIdx ON SALES.INDEXTEST(SpecialOfferID) INCLUDE (SalesOrderID,UnitPrice,OrderQty)

SELECT

SalesOrderID

,**SUM**(UnitPrice*OrderQty) AS TOTAL
FROM SALES.INDEXTEST **WITH(INDEX(NCIdx))**

WHERE SpecialOfferID=2

GROUP BY SalesOrderID

ORDER BY TOTAL;

--17 pages,non clustered index seek, sort for grouping .
--now since i group on saleseorderid, if i include that in the sort order, it would become faster

DROP INDEX CluIdx ON SALES.INDEXTEST;
DROP INDEX NCIdx ON SALES.INDEXTEST;
CREATE CLUSTERED INDEX CluIdx ON SALES.INDEXTEST (SALESORDERID);
CREATE NONCLUSTERED INDEX NCIdx ON SALES.INDEXTEST(SpecialOfferID,SalesOrderID) INCLUDE (UnitPrice,OrderQty)

SELECT

SalesOrderID

,**SUM**(UnitPrice*OrderQty) AS TOTAL
FROM SALES.INDEXTEST **WITH(INDEX(NCIdx))**

WHERE SpecialOfferID=2

GROUP BY SalesOrderID

ORDER BY TOTAL;

--17 pages,non clustered index seek,no sort for groupingthe sort operation has now disappeared

--THESE QUERIES ARE SAME

```
SELECT  
Color  
FROM Production.Product  
WHERE Color IS NOT NULL
```

```
GROUP BY Color;
```

--THESE QUERIES ARE SAME ARE NOT SAME!!

```
SELECT  
COUNT(DISTINCT Color)  
FROM Production.Product  
WHERE Color IS NOT NULL
```

```
GROUP BY Color;
```

```
SELECT  
COUNT(Color)  
FROM Production.Product  
WHERE Color IS NOT NULL
```

- Table Variables (@) are not performant as statistics about them are maintained but temp variables do maintain that & thus give good performance as a good query plan can be selected.
- Using same CTE multiple times in the same query might have performance repercussions. Think about it.
- Using performance repercussions of temp table. equijoin of result in table var & sorted inputs both sides = Merge JOIN. Nested loop join & sorted inputs both sides = Hash JOIN. Nested rowsets? None of the inputs sorted (indexed) for smaller rowsets? Equi join or Non-equi joins for index).
- JOIN for equi or Non-equi joins are embedded in index. JOIN for equi scans (some order as of rows inserted into the table). Collected Allocation ORDER SCAN
- Physical order scans (some order as of rows inserted into the table) Table INDEX ORDER SCAN
- Collected INDEX ORDER SCAN
- optimizer hints (table, query & join hints) are directives, not hints.

- Turn the setting to use double quotes as identifier delimiter.
Evict SP help tent (or) Select OBJECT_DEF

Select OBJECT_DEFINITION (OBJECT_ID)

Differences between 'ORDER BY' used for filtering ('idlo_sales'):

or win dosing purposes vs ORDER BY 'Dose' or 'Presentation ID'

Running totals
in ~~table~~: possible: ~~method~~ is temp. table
of running sum
local parameter

→ @mytbl is the name given to a global temp table. @@@ system func

• PK & FK inner join gains performance if we also declare a non-clustering key:

check constraint evaluates to 'TRUE' with NULL values (ignoring the NOT NULL constraint).

↳ causes expressions to evaluate to **NOT TRUE** when one explicitly
'UPF' vs 'If' → how does DEFAULT constraint behave when one explicitly
specifies a non-value if specified.

If you specify a NULL value vs. isnull, no value = / 0, then expression, use @

To turn values into a table info.)
class **VALUES** triple Value Constructor. **INSERT** identity

Even with 'IDENTITY INSERT ON', you can only 'UPDATE' u need Sequence

values), not replace them; run 'SEPARATABLE' to prevent infection.

Objects in table must hold lock or serial checking for value before reading through double

issues arising about using single quotes instead of double quotes; reason is that SQL should be used instead of identifiers. Reason is I think quotes

[] show, ANSI Standard box items can turn off the use of word processing features.

The server has an option now to send *Digest-Authenticate* and *Proxy-Authenticate* responses.

The `for` clause can be omitted if the `BEGIN-END` block not required in stored procedure.

→ To SET the Variable TO, shown in round brackets, can also use:

Subquery ^{since} query (for $\forall x = y$) inside using ; = inside operations. (mod views) 8 pages

→ UDF can't do DDL & DML (table, index) of own
→ in "icook": 8KB page, is file-based index. RID vs KEY

Scan vs. search -
- Leptent: Heap vs. Balanced trees (AVL tree). Nonclustered indexes on a clustered

— ~~non-clustered~~ non-clustered lookup (on balanced tree) uses KEY lookup.

indexed table. NULLs but waste. How to allow multiple NULLs in each page of the non-clustered

You can accomodate more rows in the clustered-index table (balanced tree) than in the clustered-index table (shorter than Rij).

When you do a query lookup, because clustering key (key lookup) is part of the clustered index, as it is the ~~table~~ clustered index, it is the type of index.

→ There can be only one root node in a balanced tree. (So when do we get it?)

→ used by Primary key or unique key → GROUP BY & ORDER BY